

Frozen Yet Adaptive: Single-Episode Repair of Vision–Language–Action Policies from Proprioceptive Residuals

Anonymous authors

Paper under double-blind review

Abstract

A vision–language–action (VLA) policy is a large frozen network whose training assumed a healthy robot. When the execution hardware drifts—a miscalibrated tool frame, a stiff or weakened joint—the policy is silently wrong, and fine-tuning it is out of reach for most deployments. We show that the fault can be identified and cancelled within a single episode, from the robot’s own motion, by adapting six numbers at the policy’s action interface: a plant model and a sensitivity matrix identified once on healthy data turn the proprioceptive residual into an estimate of the fault, and a projected, normalised gradient law drives it to zero. No gradients through the policy, no reward, no reference trajectory, and no learning across episodes. On four LIBERO suites with a frozen $\pi_{0.5}$, paired episodes go from 23% to 65% success (exact McNemar $p = 6.1 \times 10^{-13}$, $n = 120$), with every suite individually significant; across 690 paired episodes spanning additive, structured, time-varying and multiplicative faults the correction fixes 275 episodes and breaks 7. The same law repairs two further backbones from two other developers (OpenVLA-OFT, NVIDIA GR00T N1.7) with the same calibration, and, in joint space with a calibration of their own from a few healthy rollouts, two further manipulators: a bimanual ALOHA and a humanoid Fourier GR1 under GR00T N1.5: on two tasks, 3% $\rightarrow$ 63% and 3% $\rightarrow$ 43%, each its healthy rate, zero regressions in sixty paired episodes). A tuned integral baseline without the plant model reaches 35% where the method reaches 95%. We state the two conditions the evidence supports—identifiability, and a reparability criterion comparing the task’s spatial margin against the correction’s variation—and the failure modes that bound the method, including faults invisible to proprioception and a joint lock that is identified but cannot be cancelled from the action interface.

1 Introduction

Vision–language–action models map images and a language instruction to robot actions and are trained on demonstrations collected from a particular robot in a particular state of repair. Deployed, they are frozen. The robot is not: tool frames drift, calibrations go stale, joints lose effectiveness. Each changes the map from the policy’s action to the robot’s motion, and the policy, which never compares commanded motion with achieved motion, has no mechanism to notice.

The standard remedies are to fine-tune on faulted data, which requires labelled faults and the compute to update billions of parameters, or to learn a residual correction offline against reference rollouts. Both need the fault to be represented in a dataset before it can be corrected—the closest contemporary method, J-PARC (Jo et al., 2026), trains its calibrator per environment from faulted rollouts against fault-free references and teacher targets. Neither is available to a robot that degrades while running.

This paper takes the classical control view. A fault at the action interface is a disturbance on a plant whose healthy dynamics can be identified once; the disturbance is observable in the difference between the motion the healthy plant would have produced and the motion actually measured; and an adaptive law with a handful of parameters can cancel it online.

The question is whether this survives a 3.35B-parameter policy, contact-rich tasks, and a residual that carries model error as well as the fault. Our contributions:

1. A single-episode repair loop for frozen VLAs (§4) needing only a plant model and a sensitivity matrix identified on healthy data, both properties of the robot rather than of the policy or the task. The estimate resets every episode and reconverges within a few tens of control steps.
2. Evidence at scale, paired (§6): four suites, three backbones, three manipulators, four fault families, three severities and four time profiles, evaluated episode-for-episode against the frozen policy with exact McNemar tests, a healthy control on every cell, and 750 paired episodes in aggregate.
3. Two conditions, stated as propositions (§5): a constant input fault and a constant output bias are not separable within one episode, so the healthy bias must be measured outside it; and repair succeeds when the correction’s within-episode variation fits inside the task’s spatial margin, both quantities being measurable before any repair attempt. The second predicts every outcome on two further manipulators, including the ones continuous adaptation fails.
4. The boundary (§7): what the method cannot see (sensor and camera faults), cannot fix (a policy that fails the task unfaulted), and gets wrong (one healthy episode in twenty on one backbone, caused by a sustained phantom).

2 Related work

VLA policies and their fragility. π_0 and $\pi_{0.5}$ generate action chunks by flow matching from a vision–language backbone; OpenVLA and its OFT variant use an autoregressive backbone with a regression head. Both are evaluated on LIBERO in a fixed simulated robot. Robustness to execution faults is not part of the evaluation protocol of either.

Residual calibration of frozen VLAs. The closest work is J-PARC (Jo et al., 2026): joint-level physical faults (lock, limited range, friction) on frozen $\pi_{0.5}$ and OpenVLA-OFT across the same four LIBERO suites, repaired by a residual calibrator trained offline, per environment, from faulted rollouts with fault-free references and teacher targets (+6.8 and +4.5 points under joint lock, unpaired). We use no faulted data, no references and no teacher: the plant is identified once on healthy motion and the fault online within the episode it occurs in, on paired episodes. Their faults enter below the controller; ours enter at the action interface, and §6.5 tests their class with our law. Closed-loop affine activation editing (Das et al., 2026) steers a frozen policy by editing internal activations against an observed objective; our correction lives outside the network and needs no access to activations.

Adaptive and fault-tolerant control. Our law is a disturbance observer with feedforward cancellation, in the MRAC gradient form $\dot{\theta} = -\Gamma\phi e$, with the standard robustness modifications (deadzone, normalisation, projection). What is new is not the law but its target: a frozen generative policy whose command distribution determines what is identifiable, and a task whose spatial margin determines what is repairable.

3 Setting

A frozen policy π_ω emits, at 20 Hz (LIBERO, GR1) or 50 Hz (ALOHA), an action chunk from which the controller consumes $a_t \in \mathbb{R}^d$ per step. For LIBERO a_t is a six-dimensional Cartesian increment plus gripper, executed by an operational-space controller; for ALOHA it is fourteen and for the GR1 humanoid twenty-nine absolute joint targets executed by position servos. The measured state y_t is the end-effector increment (LIBERO) or the joint position (ALOHA, GR1).

A fault corrupts the executed action:

$$a_t^{\text{exec}} = a_t + f \quad (\text{additive}), \quad a_t^{\text{exec}} = \text{diag}(g) a_t \quad (\text{multiplicative}),$$

constant or varying in time. The fault is injected between the policy and the controller. For LIBERO this models a miscalibrated tool frame or a stale hand–eye calibration; it is not a

joint-level actuator fault, which would enter below the controller where the fault-to-motion map is state dependent (§8).

The policy never sees f . It sees images and the proprioceptive state, and $\pi_{0.5}$ -LIBERO in particular ignores the state input entirely (zeroing it changes the action by exactly 0.00000), so a perception-side fault cannot be constructed for it and an action-side fault cannot be corrected by it.

4 Method

4.1 Plant, residual, and sensitivity

On healthy rollouts we identify a per-channel FIR plant

$$\hat{y}_t = \sum_{\ell=0}^K h_{\ell} u_{t-\ell} + c, \quad K = 6, \quad (1)$$

where u_t is the command we believe we sent (the policy’s action plus any correction we applied). For LIBERO the plant predicts the increment; for ALOHA it predicts the position, because there the command is an absolute target that a servo tracks and a target offset lives in the steady-state position error, not in the increment (§6.7). The residual

$$r_t = y_t - \hat{y}_t \quad (2)$$

is what the healthy plant cannot explain. Because u_t includes the correction and the world executes $u_t + f$, $r_t \approx Mf$ independently of the correction, so the residual estimates the total fault. $M = \partial(\text{motion})/\partial(\text{fault})$ is a $d \times d$ sensitivity matrix measured once by open-loop replay: a recorded healthy command sequence is replayed with and without a per-axis probe, and the mean motion difference per unit probe gives each column. On the LIBERO plant M is diagonal-dominant with $\text{cond}(M) = 3.0$ and is magnitude independent to within 14% between probes of 0.02 and 0.05; on ALOHA $M \approx I$ with $\text{cond}(M) = 7.3$; on the GR1 the arm block is $0.99 I$ by a direct step response from a held pose (the replay probe is contaminated on an arm that touches objects during the replay).

4.2 The adaptive law

Per control step,

$$e_t = M^{-1}r_t - b, \quad \tilde{e}_t = \frac{e_t \mathbb{I}[\|r_t\| > \delta]}{1 + \|r_t\|^2/\rho^2}, \quad \hat{f}_{t+1} = \Pi_{[-\kappa, \kappa]}(\hat{f}_t + \gamma(\tilde{e}_t - \hat{f}_t)), \quad (3)$$

and the next action is sent as $a_t - \hat{f}_t \odot m$, where m selects the channels to correct. The terms are the standard robustness modifications and each earned its place by fixing an observed failure: the deadzone δ stops plant noise integrating into drift, the normalisation damps outlier steps (contacts, joint limits), the projection κ keeps the estimate in a physical box, and b is the healthy-plant phantom—the value $\hat{f}$ settles at with no fault present—measured offline and subtracted, exactly as a sensor is zeroed before it is trusted (Proposition 1 says why it cannot be fitted within the episode).

Remark 1 (Constants are ratios of measured scales). Every threshold in equation 3 is compared against a quantity with a scale: δ against the residual norm of the healthy plant, ρ against the residual norm under the fault, κ against the expected fault, the excitation gate of §4.4 against the regressor’s magnitude. Six silent failures in this work came from carrying a constant tuned on one plant to another—a deadzone larger than the residual it gated, a normaliser that attenuated every ALOHA update to 6%, a normaliser fed the residual of joints the law was not estimating, a projection bound that coincided with the true fault and was read as a measurement. The rule that follows is mechanical: print the scale, then set the constant as a multiple of it, on every new plant.

The norms in equation 3 are taken over the channels being corrected: on the humanoid the twelve unmodelled hand joints carry three times the arm’s residual, and a norm over

all joints attenuated every arm update until the estimate settled at 13% of the fault. The fixed point of equation 3 is attenuated by $1/(1 + \|Mf\|^2/\rho^2)$; at the measured $\|r\|$ on LIBERO this is 5%, matching the observed 2% shortfall, and an innovation-normalised variant ($e_t = M^{-1}(r_t - M\hat{f}_t)$, unbiased fixed point) was indistinguishable there (0.12σ) at twice the variance. We keep equation 3 on LIBERO and ALOHA and use the innovation form on the humanoid, where the fault is three times the healthy residual and the bias would be 50%.

4.3 Which channels to correct

The separation test—the estimate under a fault minus the estimate on matched healthy episodes—is the only evidence of identification we accept; a raw $\hat{f}$ carries the plant phantom and cannot distinguish identification from bias. For the uniform six-axis fault used across the LIBERO suites, rotation identifies (+0.049 against a true 0.050) and translation does not (sign-wrong), so the headline correction is restricted to rotation. That restriction is a property of that fault, not of the channel: on a translation-only fault the same law identifies translation at 79%–106% and repairs it (§6.3). The uniform fault differs because single-axis responses stop superposing at ± 0.05 , the signature of a contact or joint limit reached only when all six axes move.

4.4 Multiplicative faults

For a gain fault the regressor follows from the plant rather than being assumed. The world executes gu and the FIR responds to what it is given, so per channel

$$y_i - \hat{y}_i = \beta_i \phi_i, \quad \phi_i = \hat{y}_i - c_i, \quad \beta_i = g_i - 1, \quad (4)$$

the regressor being the FIR-weighted command history, and the regression lives in motion units with no M^{-1} . Estimating β_i by recursive least squares on channels whose $|\phi_i|$ clears an excitation gate, the correction is the exact certainty-equivalent inverse $a_i/(1 + \hat{\beta}_i)$ with $\hat{g}_i$ projected away from zero. Two facts matter here. The compensator inverts the estimate, so estimation error is amplified by $1/\hat{g}^2$ —fourfold at $g = 0.5$ —and accuracy sufficient for an additive fault is not sufficient for a multiplicative one. And a wrong regressor is not repaired by adding parameters: regressing on the instantaneous command instead of ϕ_i produced a 33% phantom loss of effectiveness on a healthy robot, an intercept made it worse, and the derived regressor shrank the phantom ninefold and took the x -channel separation from 1.4σ to 18σ .

5 Two conditions

Proposition 1 (A constant input fault and a constant output bias are not separable). Let $y_t = M(a_t + f) + b + \nu_t$ with M square and full rank, and f, b constant over the episode. For any b' , the fault $f' = f + M^{-1}(b - b')$ reproduces every observation exactly. Only $Mf + b$ is identified.

Proof. $M(a_t + f') + b' = Ma_t + Mf + (b - b') + b' = M(a_t + f) + b$. $\square$

Consequently the healthy phantom b must come from outside the episode (we measure it on fault-free rollouts), and a fault that switches on mid-episode is better posed than one present from the start: b is identified before onset and f after.

Proposition 2 (Repairability). Let μ be the task’s spatial margin—the largest steady residual offset at the end-effector under which the healthy policy still succeeds at its healthy rate—and let σ be the within-episode variation of the applied correction. Repair recovers the healthy rate when the residual stays within μ for the whole episode, for which $\sigma < \mu$ is necessary. Both μ (by a static-correction sweep) and σ (from the estimator’s trajectory on one faulted episode) are measurable without knowing the fault and without a repair trial.




Figure 1: The fault estimate within one episode, from zero, on three robots. Left: LIBERO spatial, $\pi_{0.5}$, uniform +0.05 fault, rotation channels (median and interquartile range over 20 episodes): two reach 0.04 by step 60, the third 0.024, all still rising; at episode end the separation test puts them at 0.047–0.049. Middle: ALOHA transfer-cube, π_0 , +0.02 rad on the left arm’s six joints; the identification episode settles at 92–96% of the offset, then held (§6.7). Right: GR1 humanoid, GR00T N1.5, +0.10 rad on the seven right-arm joints; 90–112% within 100 steps, then held (§6.8).

This is an empirical proposition, not a theorem: it states the quantity the data turned out to depend on. It predicts LIBERO (a reach with centimetres of margin; σ of 0.019 cm per step; repair), ALOHA with continuous adaptation ($\mu < 0.5$ cm; $\sigma = 0.43$ cm; no repair), ALOHA with the held estimate ($\sigma = 0$; repair) (§6.7), and the GR1 humanoid both ways on two tasks (§6.8).

6 Experiments

6.1 Protocol

Paired design. Both arms of every cell run the same (task, initial state, seed) episodes and differ only in whether the correction is applied; Figure 1 shows what the corrected arm does inside one episode. We report exact McNemar tests on the discordant pairs; unpaired tests on these data are conservative and, worse, hide the regression count. Controls. Every cell has a healthy-frozen control on the same episodes: a null against a ceiling (frozen already near 100%) or a floor (frozen near 0 unfaulted) is not a null, and four cells in this work were avoidable with that control run first. Noise. $\pi_{0.5}$ samples its action flow; the identical condition scored 13/20, 13/20, 15/20 on three runs, so ± 10 points at $n = 20$ is free variation and we quote intervals. Calibration. Plant and M were identified once on libero_spatial healthy rollouts with $\pi_{0.5}$ driving and are used unchanged for every LIBERO suite and for OpenVLA-OFT; ALOHA has its own from eight healthy rollouts, the GR1 from six.

6.2 Four suites

Uniform +0.05 offset on all six action dimensions, rotation-only correction (§4.3), $\pi_{0.5}$ -LIBERO.

suite	n	frozen	corrected	fixed	broken	exact McNemar
libero_spatial	20	8/20 (40%)	18/20 (90%)	10	0	0.0020
libero_goal	40	15/40 (38%)	29/40 (72%)	15	1	0.00052
libero_object	20	5/20 (25%)	16/20 (80%)	11	0	0.00098
libero_10	40	0/40 (0%)	15/40 (38%)	15	0	6.1×10^{-5}
pooled	120	28/120 (23%)	78/120 (65%)	51	1	6.1×10^{-13}

Every suite is individually significant. The long-horizon suite goes from a frozen policy that never succeeds in forty episodes to fifteen recoveries. The single regression is a goal episode




Figure 2: The recoverability map: frozen versus corrected success on LIBERO spatial, $n = 20$ paired episodes per cell, correction on the channels each family’s separation test supports (rotation for the rotation and uniform families, translation for translation). No cell has a regression. Single-family faults degrade gracefully with magnitude; the uniform fault collapses at 0.15 because its uncorrected translation component drives the arm into contacts and limits (§7). Exact McNemar p -values are in the table below.

the frozen policy itself solves three times in four, with an ordinary-sized correction. libero_90, on which the checkpoint was never fine-tuned, scores 6/20 unfaulted and is reported as a competence floor of the policy, not a repair result.

6.3 Fault families and severity

Figure 2 and the table give the fault-family by magnitude map (two cells at $n = 40$ in Appendix D; constants ablation in Appendix A).

fault (libero_spatial, $n = 20$)	0.05	0.10	0.15
translation	18/20 → 18/20 (ceiling)	13/20 → 19/20	4/20 → 19/20
rotation	13/20 → 18/20	0/20 → 17/20	0/20 → 9/20
uniform 6-axis	8/20 → 18/20	2/20 → 11/20	0/20 → 2/20

Zero regressions in every cell. Rotation is where this policy breaks—a rotation fault at 0.10 takes it to exactly zero where translation leaves 13/20—and the correction recovers 17/20 from that floor. Single-family faults degrade gracefully with severity; the uniform fault collapses, and the comparison of the last column isolates why: the rotation fault and the uniform fault are corrected identically, and adding an uncorrected translation component drops recovery from 45% to 10%. The boundary is the plant leaving the regime where its motion is informative, not the estimator. A structured mixed-sign fault $[0, 0, 0, +.06, -.06, +.03]$ the law was never tuned on goes $2/12 \rightarrow 10/12$ with separation matching truth on every axis.

Time-varying faults (amplitude 0.10): a ramp to full over 60 steps goes $12/20 \rightarrow 18/20$ ($p = 0.031$); a non-zero-mean oscillation on rotation $13/20 \rightarrow 20/20$ ($p = 0.016$); an intermittent fault $12/20 \rightarrow 19/20$ ($p = 0.039$). The estimator tracks the waveform poorly—45–79% RMS error, not attributable to lag—and it does not matter: the correction must be roughly right in direction and scale, not a good tracker, which is why six numbers updated on CPU suffice. (A zero-mean oscillation averages away and does no damage; that cell proved nothing and is not counted.)

Loss of effectiveness (§4.4): at $g = 0.30$, $1/20 \rightarrow 17/20$ ($p = 3.1 \times 10^{-5}$) with $\hat{\beta} = (-0.723, -0.699, -0.717)$ against -0.700 ; at $g = 0.20$, $0/20 \rightarrow 17/20$ ($p = 1.5 \times 10^{-5}$). At $g = 0.5$ the frozen policy still scores 15/20 and the cell is a ceiling. Rotation gain faults are unidentifiable for this policy: no recorded step clears the excitation gate on r_x or r_z at any threshold, because the policy does not rotate the wrist. That is a property of the policy’s behaviour, not of the estimator.

6.4 Baselines

Frozen and a static oracle correction equal to the true fault bound the problem. The informative baseline is the one that asks whether the plant model and M earn their place: integral action on the raw motion error, $\hat{f} \leftarrow \hat{f} + k_i(y_t - a_t)$, swept over five gains across two orders of magnitude. On the 0.15 translation fault its best setting ($k_i = 0.005$) reaches 7/20 (35%) from a frozen 3/20; higher gains fall to 0/20–1/20; ours reaches 19/20 (95%) on the same cell. The baseline implicitly assumes achieved motion equals commanded action; the measured plant DC gain is 0.23, so with no fault at all its error is $-0.77a$ —at a command of 0.5 a phantom of 0.384 against a real fault of 0.150, $2.5\times$ the target and opposite in sign. It must creep to stay stable and saturates when it moves. Predicting $y \approx 0.23a$ instead of $y \approx a$ is the whole difference.

6.5 Faults below the controller

Every fault above enters at the action interface. J-PARC’s enter in the actuator, where the fault-to-motion map depends on configuration through the Jacobian. We inject that class in the MuJoCo model of the LIBERO Panda, below the operational-space controller: a constant elbow torque (a payload or gravity-compensation error; 5 N m leaks through the controller as an 8–9 cm drift over 40 steps, commanded or not), added elbow friction, and an elbow lock (both invisible while holding, removing 3–10 cm of commanded motion). An estimate-only probe identifies each on the translation channels and nowhere else, the lock at 15–23 healthy-phantom standard deviations. Paired cells, $\pi_{0.5}$, libero_spatial, translation corrected, healthy phantom subtracted:

joint fault (elbow)	n	frozen	corrected	fixed	broken	McNemar
torque bias 5 N m	40	20/40	32/40	15	3	0.0075
friction +10	20	11/20	14/20	6	3	0.51
friction +20	20	0/20	8/20	8	0	0.0078
lock ± 0.05 rad	20	0/20	0/20	0	0	—

The torque bias is repaired, at a cost the Cartesian cells never paid: three regressions in forty. Its estimate’s across-episode spread is 0.075, four times any Cartesian fault’s, because the offset the residual reports is right for the configuration it was measured in and wrong for the one the arm moves into, and a tracking law follows that with a lag. Heavy friction, which zeros the policy, is repaired to 40% with no regression; milder friction is not a claim ($p = 0.51$): friction acts only while the joint moves, so an estimate that was right during a reach pushes the wrong way when the arm pauses to grasp. The lock is identified, corrected and not repaired: the estimate converges within 30 steps on every episode, reading the same loss-of-motion signature as the friction cell, and the correction changes nothing. The difference is physical: friction is a force the controller overcomes by commanding more, and the correction commands more; a lock removes a degree of freedom, and no command moves the arm along a direction it can no longer produce. Identifiability is necessary for repair and not sufficient; the fault must be an input the plant still responds to. J-PARC’s offline residual reports gains under a lock, consistently: a residual trained on faulted rollouts can learn a different route through the free joints, and this law cannot.

6.6 Two further backbones, unchanged calibration

OpenVLA-OFT (7B, autoregressive backbone with an L1 action head) and NVIDIA GR00T N1.7 (3B, Cosmos VLM with a flow-matching DiT head, the official LIBERO finetune) are each served behind the same interface; every script, the plant and M are the $\pi_{0.5}$ ones.

cell (libero_spatial)	frozen	corrected	fixed	broken	McNemar
healthy control	20/20	19/20	0	1	1.0
rotation 0.10	0/20	11/20	11	0	9.8×10^{-4}
translation 0.15	0/20	14/20	14	0	1.2×10^{-4}
gain 0.20	0/20	17/20	17	0	1.5×10^{-5}

Nothing in the method was $\pi_{0.5}$ -specific. OFT is more fragile (frozen 0/20 on all three faults), recovers less on rotation (55% vs 85%; its r_y fault is only 23% identified where $\pi_{0.5}$ reached 87–100%), and the healthy control lost one episode to the law itself: a sustained phantom at four to ten times the usual size, which a confidence gate evaluated offline does not remove without delaying genuine repair by one to two seconds (§7). The three headline cells also replicate on libero_object with the same calibration (0/20 $\rightarrow$ 10/20, 0/20 $\rightarrow$ 15/20, 0/20 $\rightarrow$ 19/20).

GR00T N1.7 is a third developer and a third architecture. Its healthy control is a null (18/20 both arms, rotation phantom below 0.005). With no retuning: rotation 0.10, 0/20 $\rightarrow$ **14/20** (14 fixed, 0 broken, $p = 1.2 \times 10^{-4}$); translation 0.15, 2/20 $\rightarrow$ **15/20** (13/0, $p = 2.4 \times 10^{-4}$); gain 0.20 under the multiplicative law, 0/20 $\rightarrow$ **19/20** (19/0, $p = 3.8 \times 10^{-6}$, $\hat{\beta}$ within 2% of the true -0.80). Three families, 46 repaired, none broken. The under-identified r_y is the same channel that lags on the other two backbones: it is a property of the plant and M , not of the policy. Across the three backbones on this cell, 42 of 60 paired episodes are repaired and none is broken.

6.7 A second manipulator, in joint space

Bimanual ALOHA in gym_aloha driven by π_0 : fourteen absolute joint targets at 50 Hz, transfer-cube, 300-step episodes, healthy success 25%. The plant is identified on position from eight healthy rollouts ($R^2 = 0.989$ –1.000 on all joints), $M \approx I$. A +0.05 rad offset on the six left-arm joints displaces the gripper by 5.3 cm; the frozen policy scores 0/20 at every severity tested.

Identification transfers. Separation against the healthy control: 85–91% of the fault at 0.05 rad, 94–106% at 0.02 rad, leakage ≤ 0.007 rad onto uncorrected joints. Repair does not, at first. With continuous adaptation the task returns 0/20 at every severity, warm-started or cold, with or without an early-history defect we found and fixed. The reason is quantitative:

correction (0.02 rad fault, $n = 20$)	within-episode variation	mean residual	success
none (frozen)	—	2.12 cm	0/20
adaptive, updating throughout	0.43 cm	0.17 cm	0/20
0.019 rad (95%) applied frozen	0	0.11 cm	5/20
adapt episode 0, then hold, $n = 40$	0	0.12 cm	15/40 ($p = 6.1 \times 10^{-6}$)
healthy, no fault, $n = 40$ (law running: 14/40, $p = 0.58$)	—	—	17/40

A static-correction sweep puts the task’s margin under 0.5 cm: a 90% correction held for the whole episode restores nothing, 100% restores the healthy rate. The updating law reaches a mean residual well inside that margin and still fails, because its correction moves by 0.43 cm within an episode—the whole margin—on an arm that is precise from step one. The same number, applied stationary, repairs. Identify, then hold—adapt during one sacrificial episode, hold the estimate thereafter, the cost of a calibration procedure—recovers 15/40 from a frozen zero (held estimate 0.019–0.020 rad, 95–101%), indistinguishable from the healthy arm’s 17/40 on the same seeds ($p = 0.75$), while the law on a healthy arm holds a phantom below 0.001 rad and changes nothing. This is Proposition 2, invisible on LIBERO because a Cartesian reach tolerates centimetres of jitter.

6.8 A humanoid

Fourier GR1 upper body (two 7-joint arms, two 6-joint dexterous hands, a waist; 29 absolute joint targets at 20 Hz) in NVIDIA’s RoboCasa GR1 simulator, driven zero-shot by GR00T N1.5, for which GR1 is a pretrained embodiment. Plant from six healthy episodes ($R^2 \geq 0.998$ on the arm), M by direct step response (0.99 I on the arm block; the replay probe of §6.7 is contaminated on an arm in contact), constants from the measured residual. A +0.10 rad offset on the working arm zeros the policy on both tasks. The simulator’s scene generator is reseeded per episode so both arms see identical scenes. The estimate is read in the reach window (steps 50–200, before the hand is on the plate) over three episodes and

the median held: one episode’s end reads a contact standoff, and was 0.2rad wrong on a wrist joint on two of three draws.

task, $n = 30$ paired	healthy	frozen	corrected	fixed	broken	McNemar	vs healthy
plate-to-plate	22/30	1/30	19/30	18	0	7.6×10^{-6}	$p = 0.55$
tray-to-plate	17/30	1/30	13/30	12	0	4.9×10^{-4}	$p = 0.45$

Held vectors are 93–104% and 73–126% of the fault per joint; on a healthy arm the law changes nothing (19/30 vs 18/30; 16/30 vs 15/30). Continuous adaptation reaches 8/30 on plate-to-plate, a third of the ceiling: Proposition 2 decides the scheme here as on ALOHA. Appendix C maps three fault magnitudes.

6.9 Hardware

The X2 humanoid protocol (a real actuator-authority fault in 134 logged rollouts, invisible in position, recoverable in effort) is in Appendix B; the paired repair on the robot is pending robot time and is the one result this paper does not yet contain.

7 What does not work, and what was wrong

Invisible faults. A proprioceptive bias and a wrist-camera misalignment both damage the policy and neither reaches the residual (separation 0.004 against 0.049 for an action fault): the method sees the action interface and nothing upstream of it.

A policy that cannot do the task. `libero_90` fails 14 of 20 sampled tasks unfaulted. Correcting the fault restores nothing because nothing was lost to it.

Harm on a healthy robot. One OFT episode in twenty was lost to a sustained phantom acted on. A dwell gate removes most spurious openings but not that one, and every setting delays real repair; separating it needs information the estimator does not have.

Refuted explanations, kept. We proposed that identifiability follows from healthy action statistics (refuted: translation identifies a single-axis fault at 79%); that the plant is non-linear on translation (refuted: M is magnitude independent to 14%; superposition fails instead); and that translation correction contributes nothing (refuted: at 0.15 it takes 20% $\rightarrow$ 95%). The record keeps each with the measurement that killed it.

8 Limitations

Faults below the controller are repaired only when they remain inputs the plant responds to: a joint torque bias and heavy joint friction are, a joint lock is not (§6.5), and the torque bias carries a regression rate the Cartesian cells never showed. On hardware a stiffness loss is invisible in position and recoverable in effort, so the sensor, not the law, is what changes. On tight-margin tasks the deployable scheme is to hold the estimate after one episode, a weaker claim about online adaptation than the LIBERO one, stated as such. The policy’s stochasticity gives ± 10 points at $n = 20$; we mitigated it with pairing and $n = 40$ where it mattered, not by more seeds. We give propositions, not a recovery theorem; a contraction argument on the composite map is the natural next step and requires assumptions we could not verify on contact-rich tasks. And every constant in the law is a ratio of a measured scale that must be re-measured on a new plant—six of our own failures came from not doing so.

9 Conclusion

A frozen VLA can be repaired within one episode by six numbers at its action interface, given a plant model and a sensitivity matrix identified once on healthy data; the evidence is paired and broad, the calibration transfers across suites, backbones and manipulators, the two conditions that govern success are measurable in advance, and the failures are named. The result the paper still owes is the one on a real robot.

AI use statement

Generative AI (Claude, Anthropic) was used throughout this work as a coding and writing assistant under the authors' direction: it wrote and edited the experiment scripts, the analysis code, the running project record from which this paper is built, and this manuscript's text; the authors set the research questions, chose the experiments, ran or supervised every run on their own hardware, and reviewed the outputs. All numbers in the paper are recomputed from stored per-episode outcomes by a stdlib-only script, so no reported figure rests on model-generated text. Three mechanistic explanations proposed during the work, some AI-generated, were refuted by our own measurements and are reported as such (§7). Related work was verified against the cited preprints. We take responsibility for the final content of this work, including text, claims and artifacts produced with the aid of generative AI.

Reproducibility statement

Every number above is recomputed from stored per-episode outcomes by `openpi/mcnemar.py` with nothing installed; the anonymised repository holds those outcomes, the plant and sensitivity calibrations, and every script named in the paper. Experiments need `openpi` at commit 15a9616, the `pi05_libero` checkpoint, and one GPU. The two-virtualenv split, the OpenVLA-OFT and ALOHA environments, and the three pins that were not obvious are in `SETUP.md`. The law's constants are stated as ratios of measured scales (§4), which is what made the calibration transfer across suites and backbones without retuning.

References

- Sourav Das, David Chiu, Shashank Hegde, and Gaurav S. Sukhatme. Steering multirobot behavior via closed-loop affine activation editing. arXiv preprint arXiv:2606.11489, 2026.
- Minjae Jo, Taeyoung Kwon, Jaehyun Chun, Yunho Jeong, and Taesoo Kim. Uncovering vulnerability of vision-language-action models under joint-level physical faults. arXiv preprint arXiv:2606.10501, 2026.

A Ablation of the law's constants and the calibration size

Headline cell ($\pi_{0.5}$, `libero_spatial`, uniform +0.05, rotation-only correction), $n = 20$ paired per row, one constant changed per row from the reference ($\gamma = 0.08$, $\rho = 0.15$, deadzone 0.008, plant fitted on three healthy episodes, 285 steps). Last column: rotation estimate at episode end (last-50-step mean) against a true 0.05 on each axis.

setting	frozen	corrected	fixed	broken	McNemar	$\hat{f}_{rx}, \hat{f}_{ry}, \hat{f}_{rz}$
reference	8/20	18/20	10	0	0.0020	0.044, 0.020, 0.044
$\gamma = 0.02$ ($\frac{1}{4}\times$)	8/20	17/20	9	0	0.0039	0.036, 0.016, 0.036
$\gamma = 0.32$ ($4\times$)	7/20	18/20	11	0	0.00098	0.044, 0.019, 0.044
$\rho = 0.05$ ($\frac{1}{3}\times$)	11/20	14/20	5	2	0.45	0.027, 0.011, 0.027
$\rho = 0.50$ ($3\frac{1}{3}\times$)	8/20	20/20	12	0	0.00049	0.048, 0.022, 0.049
deadzone 0	8/20	19/20	11	0	0.00098	0.044, 0.018, 0.044
deadzone 0.03 ($\approx$ residual scale)	8/20	17/20	9	0	0.0039	0.031, 0.016, 0.032
plant from 1 healthy episode (75 steps)	10/20	17/20	8	1	0.039	0.027, 0.008, 0.044
plant from 2 healthy episodes	8/20	17/20	10	1	0.012	0.045, 0.021, 0.043

The gain is flat across a $16\times$ range: it sets the length of the transient, which at a quarter of the reference still ends well inside a 220-step episode, and the final estimate is unchanged. The deadzone is flat until it reaches the residual's own scale, where it eats signal. The normaliser is the one constant that decides the outcome, in the direction the law's fixed point $f/(1+|r|^2/\rho^2)$ predicts: $\rho = 0.05$ against a measured residual of 0.034 puts the estimate at 55% of the fault and the cell becomes a null with two regressions; $\rho = 0.50$ removes the bias

and the cell reaches 20/20. The only setting that fails is the one chosen below the measured residual scale, which is the content of the “constants as ratios of measured scales” rule in §4. One healthy episode, 75 control steps, is enough for the plant: 17/20 with one rotation axis under-identified and the other two carrying the repair; two episodes recover the full estimate.

B Hardware protocol on the X2 humanoid

On logs from an X2 humanoid running a learned box-pickup policy, a real actuator-authority fault (stiffness, damping and feedforward scaled together, 0.1–1.3 across 134 rollouts) is invisible in joint position and recoverable in effort ($r = +0.84$); the plant is run-varying on load-bearing joints, so calibration must be within-run and in motion; the detectability floor is 0.5–1.3° of joint offset. The paired repair experiment on the robot is pending robot time and is the one result this paper does not yet contain.

C Humanoid recoverability map

Plate-to-plate, right-arm offset at three magnitudes, paired on identical scenes, identify over three episodes then hold (§6.8), $n = 30$ per arm:

offset	healthy	frozen	corrected	fixed	broken	McNemar	held estimate
0.05 rad	19/30	7/30	24/30	17	0	1.5×10^{-5}	90–108%
0.10 rad	22/30	1/30	19/30	18	0	7.6×10^{-6}	93–104%
0.20 rad	25/30	0/30	21/30	21	0	9.5×10^{-7}	94–99%

Every cell is at its healthy rate ($p = 0.18, 0.55, 0.29$ against the healthy arm) with no regression in 90 paired episodes; at 0.20 rad (11.5° on every joint of the working arm) the frozen policy never succeeds. The identification stays within 10% across a fourfold range of magnitude, with no sign of the superposition boundary the Panda shows at its largest uniform fault (§6.3).

D Recoverability-map cells at $n = 40$

The two cells of §6.3 that were significant only at the exact test’s floor or just above it were rerun at $n = 40$ (same configuration, initial states 45–48). Rotation 0.05: 27/40 $\rightarrow$ 39/40, 12 fixed, 0 broken, $p = 4.9 \times 10^{-4}$. Translation 0.10: 30/40 $\rightarrow$ 38/40, 9 fixed, 1 broken, $p = 0.021$. Every map cell that showed an effect at $n = 20$ is individually significant; uniform 0.15 is the superposition boundary and is not.